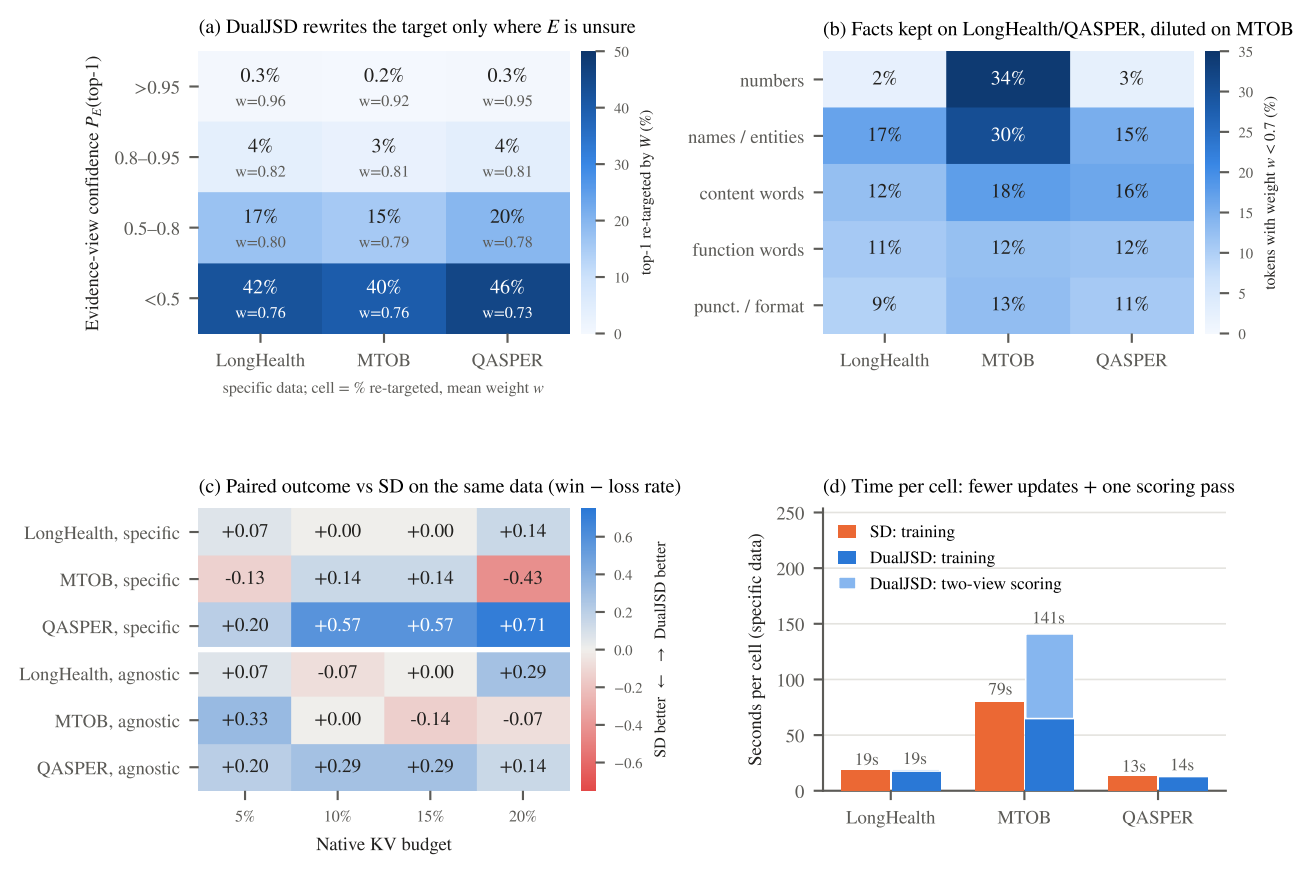
(a) DualJSD rewrites the target only where E is unsure
Evidence-view confidence P_E(top-1)
>0.95
0.8–0.95
0.5–0.8
<0.5
0.3% w=0.96
0.2% w=0.92
0.3% w=0.95
4% w=0.82
3% w=0.81
4% w=0.81
17% w=0.80
15% w=0.79
20% w=0.78
42% w=0.76
40% w=0.76
46% w=0.73
LongHealth
MTOB
QASPER
specific data; cell = % re-targeted, mean weight w
top-1 re-targeted by W (%)
50
40
30
20
10
0
(b) Facts kept on LongHealth/QASPER, diluted on MTOB
numbers
names / entities
content words
function words
punct. / format
2%
34%
3%
17%
30%
15%
12%
18%
16%
11%
12%
12%
9%
13%
11%
LongHealth
MTOB
QASPER
tokens with weight $w < 0.7$ (%)
35
30
25
20
15
10
5
0
(c) Paired outcome vs SD on the same data (win − loss rate)
LongHealth, specific
MTOB, specific
QASPER, specific
LongHealth, agnostic
MTOB, agnostic
QASPER, agnostic
+0.07
+0.00
+0.00
+0.14
-0.13
+0.14
+0.14
-0.43
+0.20
+0.57
+0.57
+0.71
+0.07
-0.07
+0.00
+0.29
+0.33
+0.00
-0.14
-0.07
+0.20
+0.29
+0.29
+0.14
5%
10%
15%
20%
Native KV budget
SD better ← → DualJSD better
0.6
0.4
0.2
0.0
−0.2
−0.4
−0.6
(d) Time per cell: fewer updates + one scoring pass
SD: training
DualJSD: training
DualJSD: two-view scoring
Seconds per cell (specific data)
250
200
150
100
50
0
19s
19s
79s
141s
13s
14s
LongHealth
MTOB
QASPER