## A. How evidence is extracted (no benchmark answers, no generation)

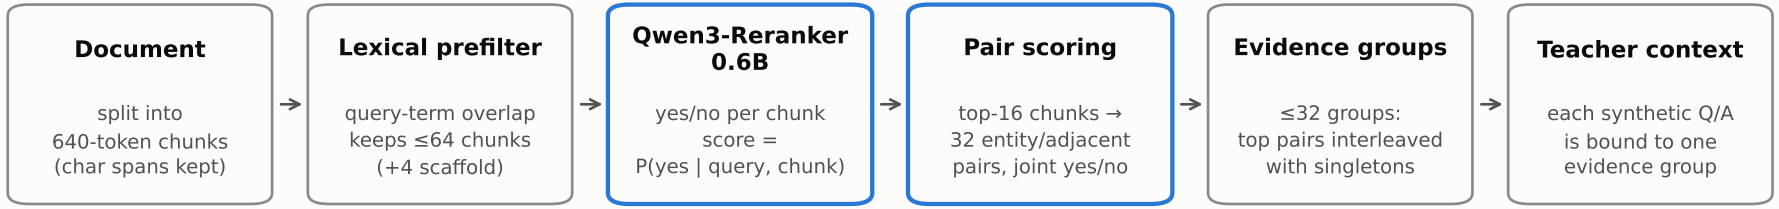


B. LongHealth: is gold evidence ranked first?
n=32 requests, 20 chunks/doc, MRR 0.72 vs 0.45

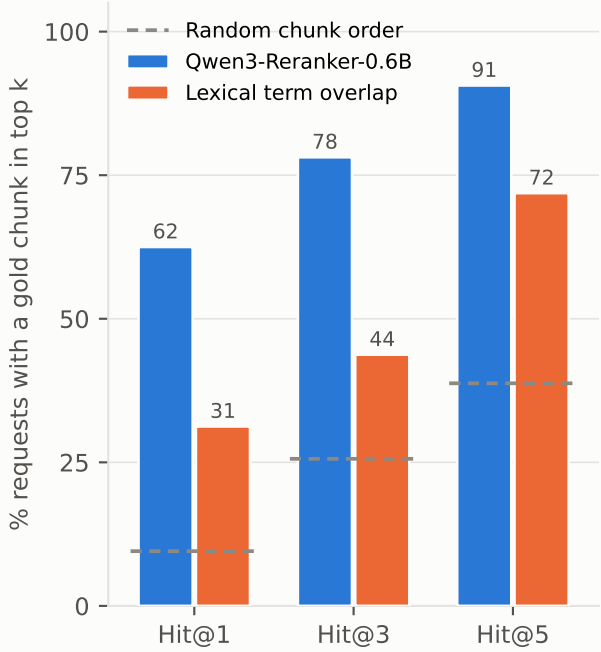


C. QASPER: is gold evidence ranked first?
n=29 requests, 9 chunks/doc, MRR 0.52 vs 0.37

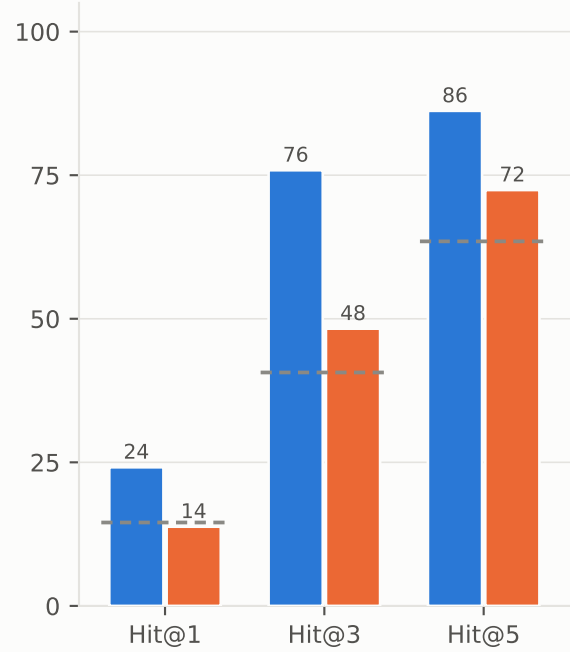


D. Does the Teacher context contain gold evidence?
(15 trajectories per request and arm)

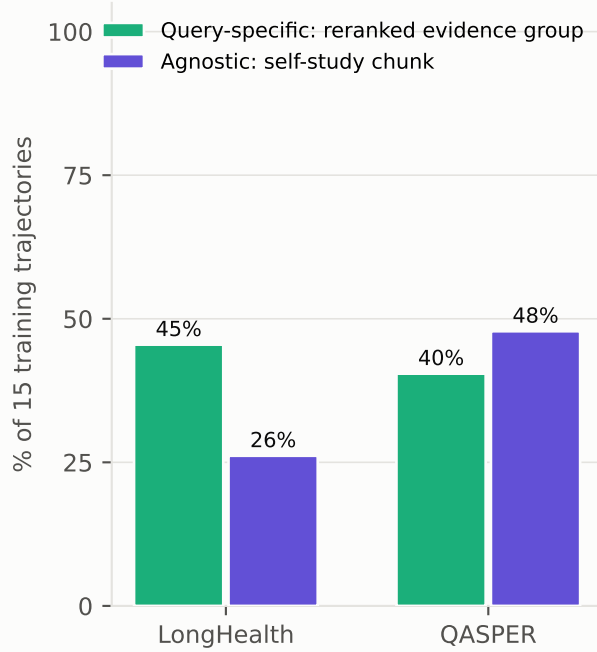


Gold = LongHealth answer_location / QASPER annotated evidence paragraphs, used only for this audit. MTOB has no evidence labels. Match = central 60-character snippet of a gold span.